

# Software Construction Laboratory

Week 6 Part 1

Lab 3

Mushi Zhou

Winter 2017

UCLA

Prof. Paul Eggert

# Some Help for Assignment 5

For Lab 6,

- Compare the performance of getchar/putchar vs read/write
- Read/write are system calls, i.e. they are unbuffered if you make these calls directly, but you can specify how many bytes to read/write with one call
- Getchar/putchar are buffered so that they read/write everything with one read/write system call to a temporary buffer, and then read/write byte by byte from that buffer

# Strace

- Strace, as we have mentioned before, is a debugging tool, that can monitor interactions between processes and Linux kernel, i.e. system calls
- Just do *strace programname* (executable, not your source file)
- By default, strace output all system calls made
- Options:
  - -e to show only selected system calls, exp: `strace -e open ls`
  - -t print time stamp of system calls
  - -r print relative time of execution
  - -c show a statistic report summary

# Some Help for Assignment 5

For HW,

- Modify your HW5 to use system calls instead
- Adding the non-case sensitivity option '-f' to comparisons
- Note: the topper() takes in unsigned characters, but the default char is signed, so you'll have to shift the value before changing to upper case
- Compare execution time of your two HWs in regards to number of comparisons
- What one do you expect to run faster?

# Outline for Week 6

- Threads vs Processes
- Multithread Performance
- POSIX Threads
- Thread Synchronizations

# What are Process and Thread

Recall from last week:

- A process (or a task) is an executing (i.e., running) instance of a program
- A thread of execution (thread) is the smallest sequence of programmed instructions that can be managed independently in a process
- In most cases a thread is a component of a process
- You can think of threads as child processes that share the parent process resources but execute independently

# Multiprocessing vs Multithreading

- These are different!
- Multiprocessing -> in hardware level, if there are more than one processor in your CPU, they are run parallelly for multiple processes
- Having multiple programs running at the same time = *multitasking*
- Multithreading -> is an execution model that allows a single process to have multiple code segments (i.e., *threads*) run concurrently within the context of that process
- Multiple threads of a single process can share the CPU in a single CPU system or (purely) run in parallel in a multiprocessing system

# Why Do We Want Multithreading?

- It is necessary!
- For example: You have a GUI application. You request a calculation that needs a fair mount of time to compute
- In a single thread mode, your GUI is going to be unresponsive until that calculation is done. You can perform nothing on the GUI during the period
- If the application is a multithreaded environment, that calculation becomes a thread, and you can perform other tasks while that calculation is running



# Multithreaded Performance

- Multithreading makes processing time much shorter!
- You will verify this in assignment 6
- How much faster depends on implementation and the natural of the problem
- Natural of the problem:
  - From embarrassingly parallel
  - To inherently serial
  - Most of problem fall in between the range

# Embarrassingly Parallel Problem

- little or no effort is needed to separate the problem into a number of parallel tasks. The sub-tasks are neutrally parallel.

- For example :  $a = ?$   $b = ?$

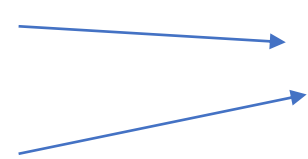
5s

$$a = c + d$$

5s

$$b = e + f$$

These can be done parallelly in  
different threads



- The execution time can be reduced from 10s to 5s by computing a and b in parallel
- In this type of problem, multithreading can drastically reduce the processing time

# Inherently Serial Problem

- For example:  $a + b$  ?

$$a = c + d$$

$$b = a + e$$

- It is not easy to parallel these since  $b$  depends on  $a$  that is been computed
- The multithreaded execution on this type of problem is not likely to yield much faster outcome

# How to Achieve Multithreading in a C program?

- POSIX Threads (Pthreads), is an execution model that allows multithreading
- It allows a program to control multiple different flows of work that overlap in time
- Each flow of work is referred to as a thread
- Creation and control over these flows is achieved by making calls to the POSIX Threads API
- When compiling the program, use `-lpthread` option to link the Pthread libraries

# Basic Steps to Use POSIX Threads

- Include <pthread.h>
- Use pthread\_create and pthread\_join calls to Pthread API
- pthread\_create (thread, attr, start routine, arg)
- This creates a new thread within a process
- The new thread execute the start routine function with the input argument arg
- The new thread has attributes specified with attr, or default attributes if attr is NULL
- If the pthread\_create() routine succeeds it will return 0 and put the new thread id into thread, otherwise an error number shall be returned indicating the error

# Basic Steps to Use POSIX Threads

- Function: `pthread_join()`
- `int pthread_join(pthread_t thread, void ** status);`
- This function suspends execution of the current process and waits for the target thread to terminate
- On a successful call `pthread_join()` will return 0
- On failure `pthread_join()` will return an error number indicating the error
- You only need to use `pthread_join` and `pthread_create` for assignment 6

# Check Error

There are returned values from the function call

For Pthread\_create:

- EAGAIN: The process lacks the resources to create another thread, or the total number of threads in a process would exceed PTHREAD\_THREADS\_MAX.
- EINVAL: *attr* is invalid.

For Pthread\_join

- EINVAL: The value specified by the thread is not a valid thread
- ESRCH: The specified thread is already detached
- EDEADLK: The join would result in a deadlock or the value of thread specifies the calling thread

# What Could Go Wrong in Multithreading

- `*a` is not defined initially
  - Thread A -> `*a = 5`
  - Thread B -> `*a = 6`
  - Thread C -> `printf("a = %d", *a)`
- 
- Assume threads a,b,c are started at the same time, what is printed?
  - There are three possibilities
  - This is called race condition!



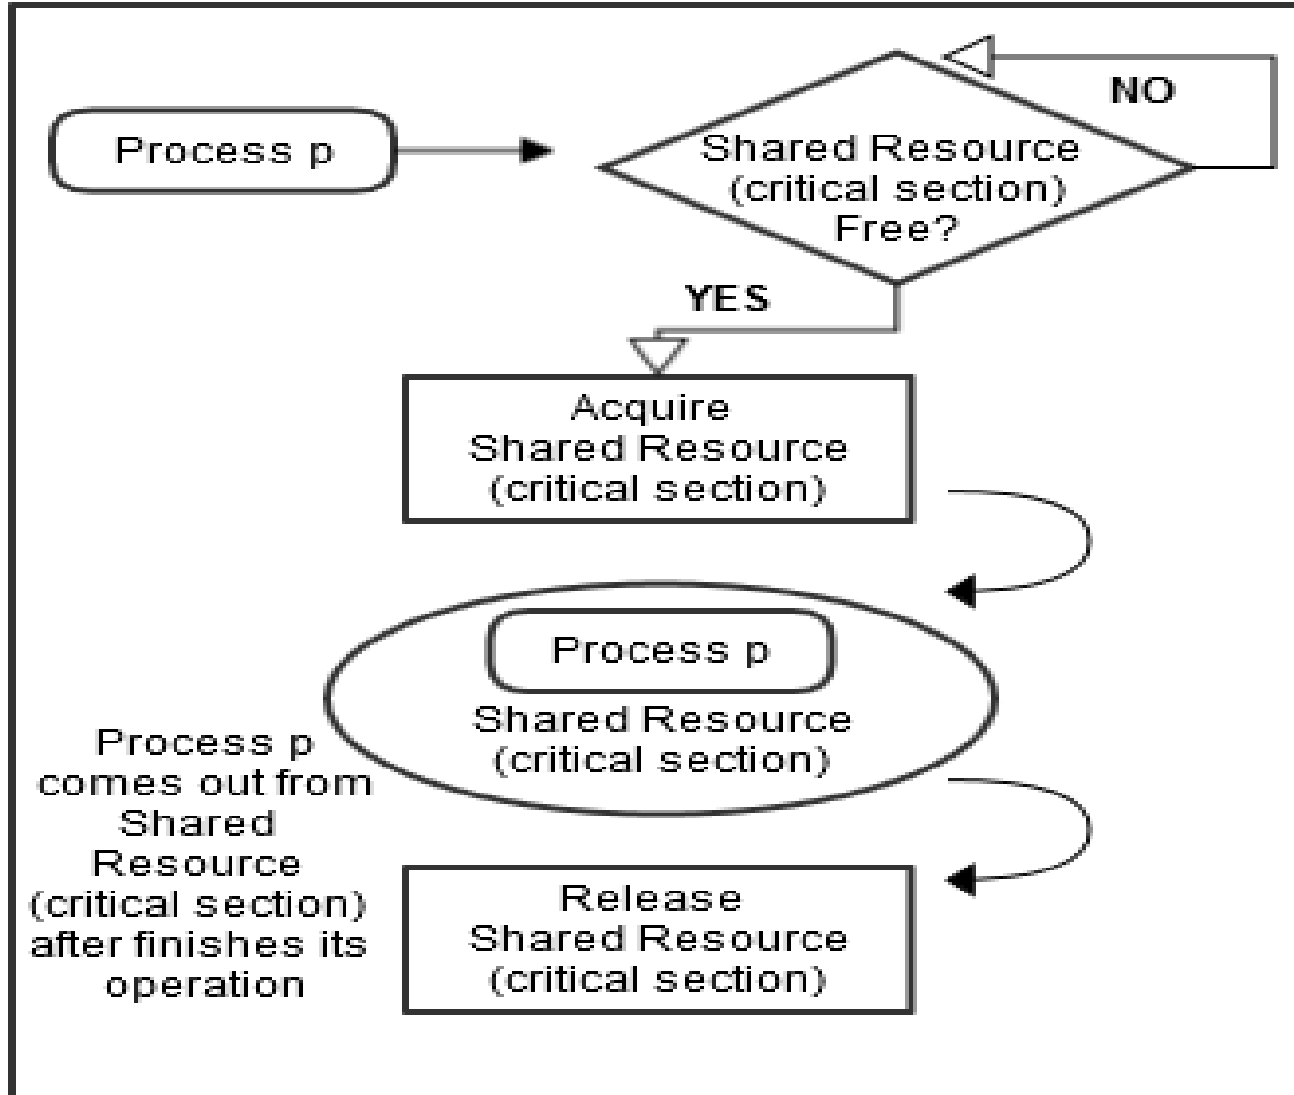
# Race Condition

- A **race condition** is an undesirable situation that occurs when a device or system attempts to perform two or more operations at the same time, but because of the nature of the device or system, the operations must be done in the proper sequence to be done correctly
- In other words, it is the behavior of an program where the output is dependent on the sequence or timing of other uncontrollable events
- So we need to deal with this kind of problem
- Synchronization

# Synchronization in Multithreading

- A mechanism which ensures that two or more concurrent threads do not simultaneously execute some particular program segment known as critical section
- Processes' access to critical section is controlled by using synchronization techniques
- When one thread starts executing the critical section (serialized segment of the program) the other thread should wait until the first thread finishes

# Illustration of Synchronization



The action of getting to the shared resource is commonly referred as “acquiring the lock” of the resource

# More on Synchronization

- Synchronization needs to be implemented carefully so that it does not lead to deadlock -> where all threads are looping forever waiting to acquire the lock to the shared resource where the lock is actually held by one of the waiting threads
- Pthread API provides synchronization by:
  - Function: `pthread_mutex_lock()`
  - Function: `pthread_mutex_trylock()`
  - Function: `pthread_mutex_unlock()`
  - Etc.

# More Resources on Multithreading

- Pthread API

<http://www.cs.wm.edu/wmpthreads.html>

- Common Synchronization Problems

[https://en.wikipedia.org/wiki/Synchronization \(computer science\)](https://en.wikipedia.org/wiki/Synchronization_(computer_science))

- Multithreading vs Multiprocessing vs Multitasking

<https://gabrieletolomei.wordpress.com/miscellanea/operating-systems/multiprogramming-multiprocessing-multitasking-multithreading/>